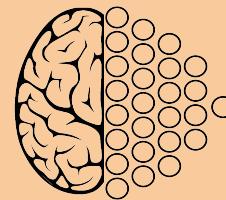


[[Video](#)]



Deep State Space Models

Karan Bania, Yash Bhisikar



S4



S5



S6

27th September 2024

Hype around SSMs



Albert Gu
@albertgu

Quadratic attention has been indispensable for information-dense modalities such as language... until now.

Announcing Mamba: a new SSM arch. that has linear-time scaling, ultra long context, and most importantly--outperforms Transformers everywhere we've tried.

With @tri_dao 1/



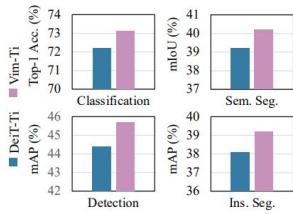
Paving the way to efficient architectures:
StripedHyena-7B, open source models offering a glimpse into a world beyond Transformers



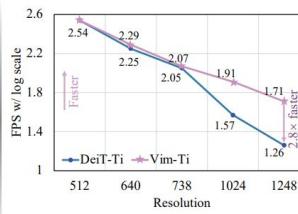
Aleksandar Botev
@botev_mg

We present Griffin: A hybrid model mixing a gated linear attention with local attention. This combination is extremely effective in leveraging the efficient benefits of linear RNNs and the expressiveness of transformers. Scaled up to 14B!

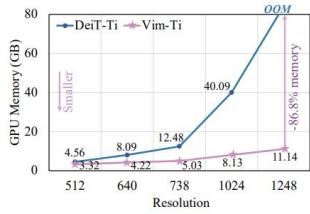
<https://arxiv.org/abs/2402.19427>



(a) Accuracy Comparison



(b) Speed Comparison



(c) GPU Memory Comparison



Mamba: <https://arxiv.org/abs/2401.09417>

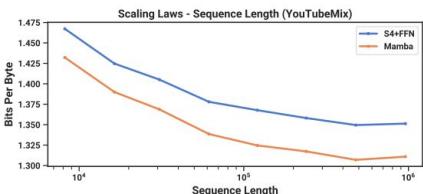
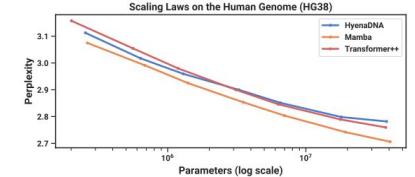


Tri Dao

With @albertgu, we're collaborating with @togethercompute and @cartesia_ai and releasing a Mamba-3B model trained on 600B tokens on the SlimPajama dataset (Mamba-3B-SlimPJ). It's among the strongest 3B models, matching the performance of strong Transformers (BTLM-3B).

1/

	Mamba-3B-SlimPJ	BTLM-3B-8K	StableLM-3B-4E1T
Number of params	2.77B	2.65B	2.80B
Number of tokens	604B	627B	4T
Training FLOPs	1.01E22	1.22E22	8.33E22
BioQ	71.0	70.0	75.5
PIQA	78.1	77.2	79.8
HellaSwag	71.0	69.8	73.9
WinoGrande	65.9	65.8	66.5
ARC-e	68.2	66.9	67.8
ARC-c	41.7	37.6	40.0
OpenBookQA	39.8	40.4	39.6
RACE-high	36.6	39.4	40.6
TruthfulQA	34.3	36.0	37.2
MMLU	26.2	28.1	44.2
Avg accuracy	53.3	53.1	55.5



Mamba on Language, DNA and audio data:
<https://arxiv.org/abs/2312.00752>

And many more...

Overview

- History: *RNNs, Transformers and Problems (Yash)*
- SSM Fundamentals & S4^[1]: *Efficiently Modeling Long Sequences with Structured State Spaces (Karan)*
- S5^[2]: *Simplified State Space Layers for Sequence Modeling (Yash, 11th Oct)*
- Mamba^[3]: *Linear Time Sequence Modeling with Selective State Spaces (Karan)*
- Event-SSM^[4]: Scalable Event-by-event processing of Neuromorphic sensory signals with Deep State-Space Models (*Yash, 11th Oct*)

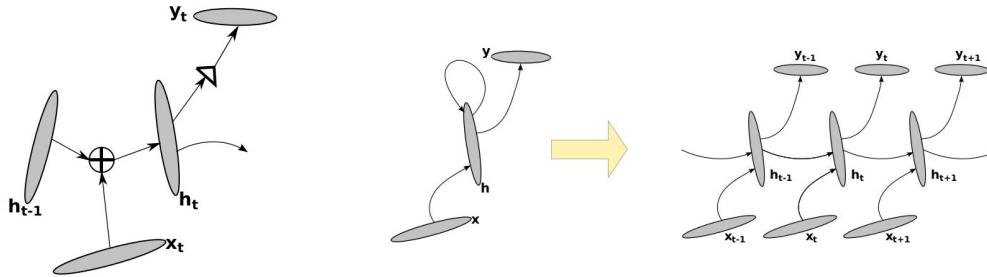
Most images have been taken from [\[1\]](#), [\[2\]](#) and [\[3\]](#) and [this](#) DeepMind talk in UCL.
(psst, We got to work with Mamba's CUDA kernels!)

Preface

- These models and many slides have a lot of **math!**
- They also involve a lot of CS fundamentals, and **out-of-the-box programming.**
- We think that regardless of Machine Learning these papers should also be treated as amazing **thought experiments**, it's definitely not the case that these models are wack at ML xD; it's just that they fully deserve all the hype around them!
- These are truly revolutionary architectures, let's start with “Is attention is all we need?”

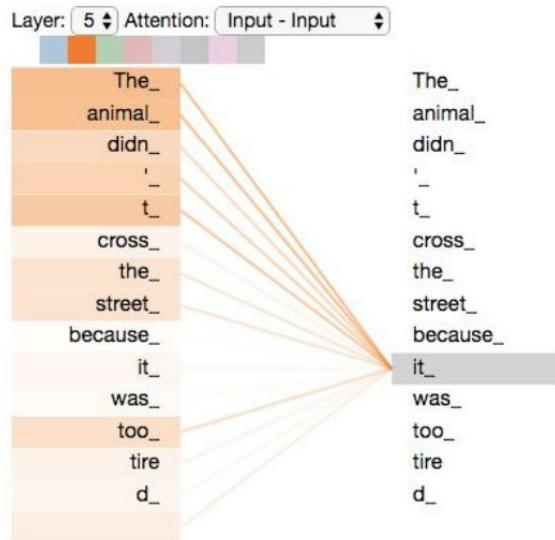
History: RNNs & LSTMs

- ~1925 Ising-Models (untrained RNNs)
- ~'72-'81 Hopfield networks (trained RNNs)
- ~'80-'90 Amari, Rumelhart et al., Werbos, etc. – Back-propagation and BPTT are introduced/popularized
- ~'90-'20s Expressivity results (Hava Siegelmann & Sontag in 1991): **RNNs are Turing-Complete**
- ~'92-'24 Bengio et al., Hochreiter & Schmidhuber: RNNs are hard to train (**vanishing/exploding gradients problems**)
- ~'01-'10 Echo State Networks / LSM as an answer to the trainability problem
- 1997 Hochreiter & Schmidhuber: LSTMs
- 2014 Graves shows LSTMs to work at “scale”
Sutskever et al. Seq2Seq Model
Chung et al. 2014, GRU



- BP was adopted (as BPTT) to train them
- Early expressivity results made RNN very desirable architecture
- Allowed to condition on an arbitrary length sequence
- Exhibits optimization issues (vanishing/exploding gradient) and scalability issues (required sequential computation)

History(?): Transformers



$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$

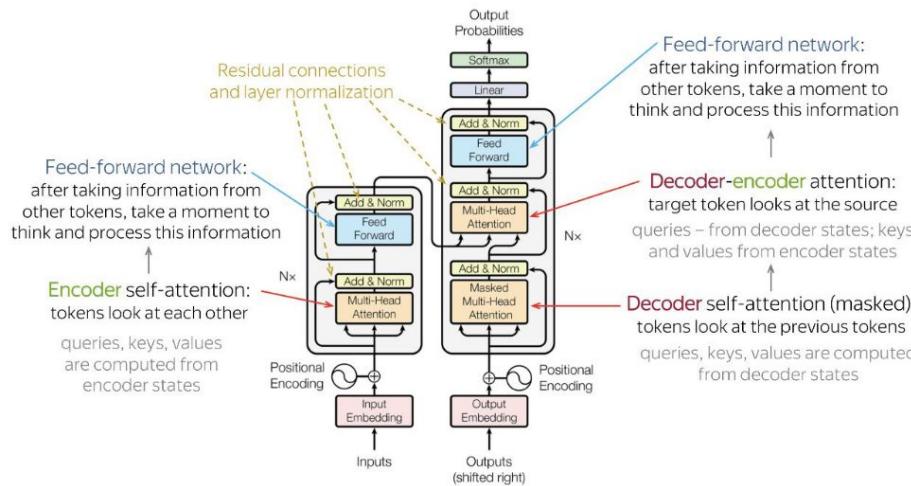
where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

$h = 8$ parallel attention layers, or heads.

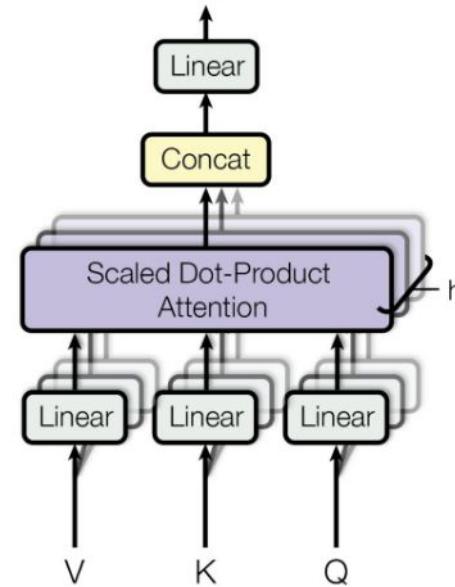
Learnable parameter matrices

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

History(?): Transformers



Multi-Head Attention



SSM Fundamentals and S4

- Before we move forward, **two** branches with “SSMs”

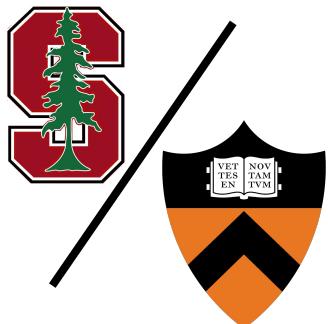


LRU



Griffin / Hawk

Much More
DL-Like.



S4



S5 / Mamba

Much More
Mathematical.



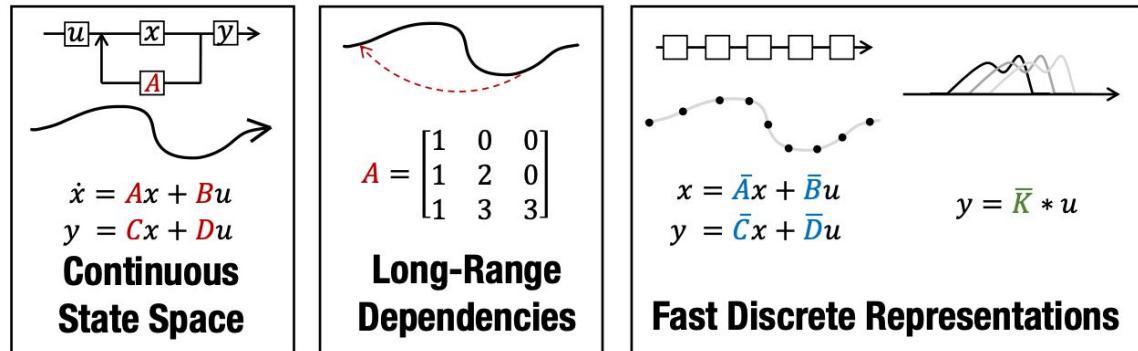
This
presentation

SSM Fundamentals and S4

- 2020 HiPPO is introduced
- 2021 Linear State Space Layer (LSSL)
- 2022 Structured State Space Model (S4)
- 2022 SaShiMi
- 2022 Hyena
- 2022 Diagonal Structured Space Models (S4D)
- 2022 Liquid State Space Models
- 2022 Simplified State Space Models (S5)
- 2022 SGConv
- 2022 Hungry Hungry Hippos (H3)
- 2023 Mega
- 2023 Linear Recurrence Units (LRU)
- 2023 RWKV
- 2023 RetNet
- 2023 2-D SSMs
- 2023 Mamba
- 2023 Vision-Mamba

Basically fixed 2 **important** problems with RNNs,

- **Stable** training,
- **Scalable** training.



SSM Fundamentals and S4

- Key ideas:
 - Continuous time interpretation,
 - Specific initialization,
 - Discretization + Diagonalization.
- Continuous time interpretation,
 - **Original** formulation of state-space models.

$$x'(t) = \mathbf{A}x(t) + \mathbf{B}u(t)$$

$$y(t) = \mathbf{C}x(t) + \mathbf{D}u(t)$$

SSM Fundamentals and S4

- Specific initialization,
 - **More theory**, this is from echo networks.

(**HiPPO Matrix**) $\mathbf{A}_{nk} = - \begin{cases} (2n + 1)^{1/2}(2k + 1)^{1/2} & \text{if } n > k \\ n + 1 & \text{if } n = k \\ 0 & \text{if } n < k \end{cases}$

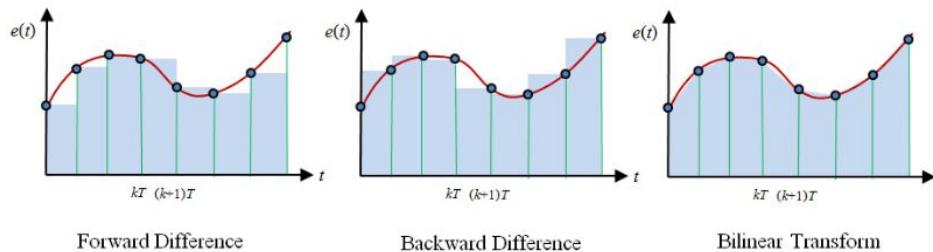
- This initialization helped get from **60% to 98%** on MNIST!

SSM Fundamentals and S4

- **Discretization** and Diagonalization.

Clearly, our input is not continuous time (speech, image-pixels, etc.), so we need to discretize the system, this lends us

$$\left. \begin{array}{l} x_k = \bar{\mathbf{A}}x_{k-1} + \bar{\mathbf{B}}u_k \\ y_k = \bar{\mathbf{C}}x_k \end{array} \right| \begin{array}{l} \bar{\mathbf{A}} = (\mathbf{I} - \Delta/2 \cdot \mathbf{A})^{-1}(\mathbf{I} + \Delta/2 \cdot \mathbf{A}) \\ \bar{\mathbf{B}} = (\mathbf{I} - \Delta/2 \cdot \mathbf{A})^{-1}\Delta\mathbf{B} \\ \bar{\mathbf{C}} = \mathbf{C}. \end{array}$$



SSM Fundamentals and S4

- Discretization and **Diagonalization**.

Why diagonalize in the first place? **Very important question!**

We will answer this but first let's see the how-to (briefly).

$$x_k = \bar{A}x_{k-1} + \bar{B}u_k$$

$$\bar{A} = PDP^{-1}$$

$$x_k = PDP^{-1}x_{k-1} + \bar{B}u_k$$

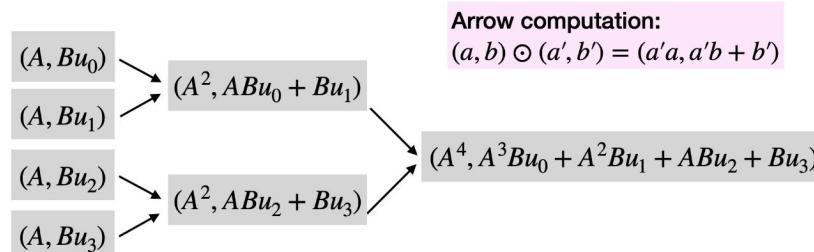
$$P^{-1}x_k = DP^{-1}x_{k-1} + P^{-1}\bar{B}u_k$$

$$\tilde{x}_k = D\tilde{x}_{k-1} + \tilde{B}u_k$$

- For arbitrary \bar{A} , D will be **complex**.

SSM Fundamentals and S4

- Now, how is this **stable**?
 - Because \bar{A} is **diagonal**, we can directly access its spectrum, and thus have **control** on recurrence blow-up. (Just parametrize such that $\lambda \leq 1$, used commonly - $\bar{A} = -e^\phi$ where Φ is **learnable**)
 - Also, recurrence is **linear**.
- How / why is it **scalable**?
 - **Associative Scans**^[2] (or interpret as convolution).



SSM Fundamentals and S4

- How / why is it **scalable**?
 - Associative Scans^[2] (or interpret as convolution).

$$x_0 = \bar{B}u_0 \quad x_1 = \bar{AB}u_0 + \bar{B}u_1 \quad x_2 = \bar{A}^2\bar{B}u_0 + \bar{AB}u_1 + \bar{B}u_2 \quad \dots$$

$$y_0 = \bar{CB}u_0 \quad y_1 = \bar{CAB}u_0 + \bar{CB}u_1 \quad y_2 = \bar{CA}^2\bar{B}u_0 + \bar{CAB}u_1 + \bar{CB}u_2 \quad \dots$$

$$\begin{aligned} y_k &= \bar{CA}^k\bar{B}u_0 + \bar{CA}^{k-1}\bar{B}u_1 + \cdots + \bar{CAB}u_{k-1} + \bar{CB}u_k \\ y &= \bar{K} * u. \end{aligned}$$

$$\bar{K} \in \mathbb{R}^L := \mathcal{K}_L(\bar{A}, \bar{B}, \bar{C}) := \left(\bar{CA}^i \bar{B} \right)_{i \in [L]} = (\bar{CB}, \bar{CAB}, \dots, \bar{CA}^{L-1} \bar{B}).$$

- +1 to diagonalization, reduces **FLOPs**, $O(H^2)$ to $O(H)$!

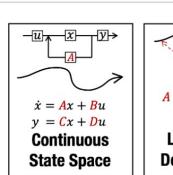
SSM Fundamentals and S4

- Then why LRU?

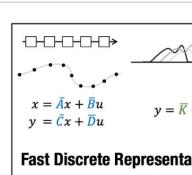
State Space Models


$$\begin{aligned}\dot{x} &= Ax + Bu \\ y &= Cx + Du\end{aligned}$$

Continuous State Space


$$A = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 2 & 0 \\ 1 & 3 & 3 \end{bmatrix}$$

Long-Range Dependencies


$$\begin{aligned}\dot{x} &= \bar{A}x + \bar{B}u \\ y &= \bar{C}x + \bar{D}u\end{aligned}$$
$$y = \bar{K} * u$$

Fast Discrete Representations

But...

- Many variants of SSMs have been proposed (S4, S4D, H3, Mamba ..)
- Multiple choices for discretization (e.g. bilinear, ZOH)
- Unclear how changes to the architecture interacts with parametrization of SSM
- Is the continuous time interpretation necessary

How can we disentangle what is important? 

Efficiently Modeling Long Sequences with Structured State Spaces

Albert Gu, Karan Goel, and Christopher Ré

Department of Computer Science, Stanford University

{albertgu,krng}@stanford.edu, chrismre@cs.stanford.edu

[[2111.00396](#)]

[[The Annotated S4](#)]

[[GitHub](#)]

SSM Fundamentals and S4

- Now the S4 paper is just connecting all the dots and making it a learnable system.
 - Continuous time interpretation ✓,
 - HiPPO initialization ✗ (almost),
 - Discretization ✓,
 - Training = Convolutional interpretation ✓,
 - Diagonalization (Motivated through computational efficiency) ✓,
 - + Extra stuff for an actual fast implementation. (Really fast!)

SSM Fundamentals and S4 [the complicated stuff]

- Actual computation and NPLR / DPLR matrices.
 - First, why Normal? Because we perform **conjugation**.
 - Second, why Low Rank? To **approximate** HiPPO matrices.
- However, this would **not** be enough, powering up a sum is still as **problematic** as powering up any other matrix => **slow** implementation.
- Naïvely, this needs $O(N^2L)$ computations to just compute the kernel, however they describe an algorithm which does the following!

Theorem 3 (S4 Convolution). *Given any step size Δ , computing the SSM convolution filter $\bar{\mathbf{K}}$ can be reduced to 4 Cauchy multiplies, requiring only $\tilde{O}(N + L)$ operations and $O(N + L)$ space.*

SSM Fundamentals and S4 [the really complicated stuff]

Algorithm 1 S4 CONVOLUTION KERNEL (SKETCH)

Input: S4 parameters $\Lambda, P, Q, B, C \in \mathbb{C}^N$ and step size Δ

Output: SSM convolution kernel $\bar{K} = \mathcal{K}_L(\bar{A}, \bar{B}, \bar{C})$ for $A = \Lambda - PQ^*$ (equation (5))

- 1: $\tilde{C} \leftarrow (\mathbf{I} - \bar{A}^L)^* \bar{C}$ ▷ Truncate SSM generating function (SSMGF) to length L
 - 2: $\begin{bmatrix} k_{00}(\omega) & k_{01}(\omega) \\ k_{10}(\omega) & k_{11}(\omega) \end{bmatrix} \leftarrow [\tilde{C} \mathbf{Q}]^* \left(\frac{2}{\Delta} \frac{1-\omega}{1+\omega} - \Lambda \right)^{-1} [\mathbf{B} \mathbf{P}]$ ▷ Black-box Cauchy kernel
 - 3: $\hat{K}(\omega) \leftarrow \frac{2}{1+\omega} [k_{00}(\omega) - k_{01}(\omega)(1 + k_{11}(\omega))^{-1} k_{10}(\omega)]$ ▷ Woodbury Identity
 - 4: $\hat{K} = \{\hat{K}(\omega) : \omega = \exp(2\pi i \frac{k}{L})\}$ ▷ Evaluate SSMGF at all roots of unity $\omega \in \Omega_L$
 - 5: $\bar{K} \leftarrow \text{iFFT}(\hat{K})$ ▷ Inverse Fourier Transform
-

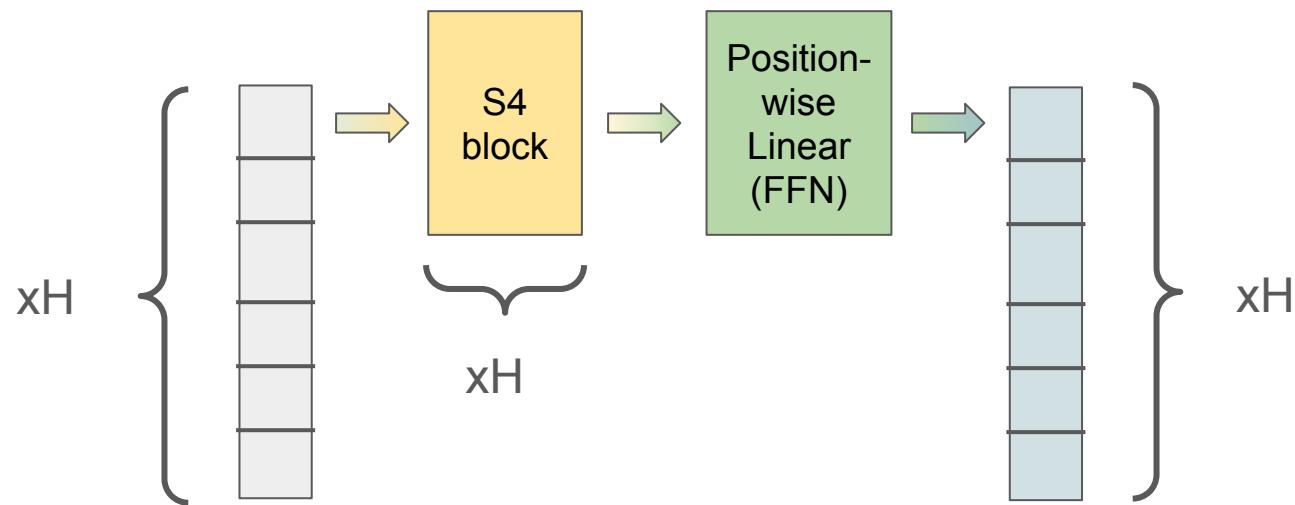
SSM Fundamentals and S4

- Key takeaways -

	Convolution ³	Recurrence	Attention	S4
Parameters	LH	H^2	H^2	H^2
Training Space	$\tilde{L}H(B + H)$	BLH^2	$B(L^2H + LH^2)$	$BH(\tilde{H} + \tilde{L}) + B\tilde{L}H$
Parallel Inference	Yes	No	Yes	Yes
	LH^2	H^2	$L^2H + H^2L$	H^2

SSM Fundamentals and S4

- Architecture (one layer) -



SSM Fundamentals and S4



Results,

Table 4: (**Long Range Arena**) (*Top*) Original Transformer variants in LRA. Full results in Appendix D.2. (*Bottom*) Other models reported in the literature. *Please read Appendix D.5 before citing this table.*

MODEL	LISTOPS	TEXT	RETRIEVAL	IMAGE	PATHFINDER	PATH-X	AVG
Transformer	36.37	64.27	57.46	42.44	71.40	✗	53.66
Reformer	<u>37.27</u>	56.10	53.40	38.07	68.50	✗	50.56
BigBird	36.05	64.02	59.29	40.83	74.87	✗	54.17
Linear Trans.	16.13	<u>65.90</u>	53.09	42.34	75.30	✗	50.46
Performer	18.01	65.40	53.82	42.77	77.05	✗	51.18
FNet	35.33	65.11	59.61	38.67	<u>77.80</u>	✗	54.42
Nyströmformer	37.15	65.52	<u>79.56</u>	41.58	70.94	✗	57.46
Luna-256	37.25	64.57	79.29	<u>47.38</u>	77.72	✗	<u>59.37</u>
S4	59.60	86.82	90.90	88.65	94.20	96.35	86.09

SSM Fundamentals and S4



Results,

Table 2: Deep SSMs: The S4 parameterization with Algorithm 1 is asymptotically more efficient than the LSSL.

Dim.	TRAINING STEP (ms)			MEMORY ALLOC. (MB)		
	128	256	512	128	256	512
LSSL	9.32	20.6	140.7	222.1	1685	13140
S4	4.77	3.07	4.75	5.3	12.6	33.5
Ratio	1.9×	6.7×	29.6×	42.0×	133×	392×

Table 3: Benchmarks vs. efficient Transformers

	LENGTH 1024		LENGTH 4096	
	Speed	Mem.	Speed	Mem.
Transformer	1×	1×	1×	1×
Performer	1.23×	<u>0.43</u> ×	3.79×	<u>0.086</u> ×
Linear Trans.	1.58 ×	0.37 ×	5.35 ×	0.067 ×
S4	1.58 ×	<u>0.43</u> ×	<u>5.19</u> ×	0.091×

SIMPLIFIED STATE SPACE LAYERS FOR SEQUENCE MODELING

!11th Oct!

Jimmy T.H. Smith^{*, 1, 2}, Andrew Warrington^{*, 2, 4}, Scott W. Linderman^{2, 3}

^{*}Equal contribution.

¹Institute for Computational and Mathematical Engineering, Stanford University.

²Wu Tsai Neurosciences Institute, Stanford University.

³Department of Statistics, Stanford University.

{ jsmith14, awarring, scott.linderman}@stanford.edu.

[[2208.04933](#)]

[[GitHub](#)]

Mamba: Linear-Time Sequence Modeling with Selective State Spaces

Albert Gu^{*}¹ and Tri Dao^{*}²

¹Machine Learning Department, Carnegie Mellon University

²Department of Computer Science, Princeton University
agu@cs.cmu.edu, tri@tridao.me

[[2312.00752](#)]

[[GitHub](#)]

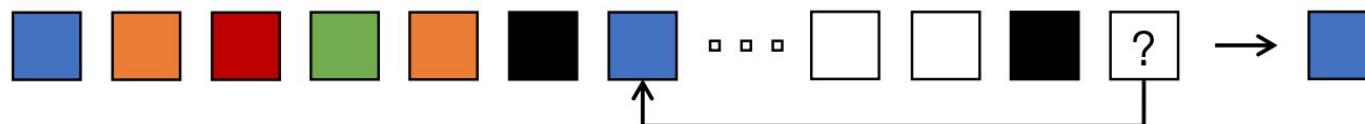
Mamba: Linear Time Sequence Modeling with Selective State Spaces

- S6 borrows a lot of parts from S4 + an important **linear-time-variant** extension.
 - Drops Complex analysis 
 - Drops HiPPO completely 
 - Linear Time Invariance 
- The paper has 3 major contributions:
 - The LTI-drop (**selection** mechanism),
 - Hardware-aware algorithm,
 - Scaling.

Mamba: Linear Time Sequence Modeling with Selective State Spaces

- Selection:
 - Why? Very closely related LSTM-gating, select data in an input-dependent manner.
 - Selectivity as the goal of language sequence modelling, effectiveness - efficiency tradeoff.

Induction Heads



Mamba: Linear Time Sequence Modeling with Selective State Spaces

- Selection:

- They make a very strong claim. It means

$$\begin{array}{l} x_k = \bar{\mathbf{A}}x_{k-1} + \bar{\mathbf{B}}u_k \\ y_k = \bar{\mathbf{C}}x_k \end{array} \quad \left| \begin{array}{l} \bar{\mathbf{A}} = (\mathbf{I} - \Delta/2 \cdot \mathbf{A})^{-1}(\mathbf{I} + \Delta/2 \cdot \mathbf{A}) \\ \bar{\mathbf{B}} = (\mathbf{I} - \Delta/2 \cdot \mathbf{A})^{-1}\Delta\mathbf{B} \\ \bar{\mathbf{C}} = \mathbf{C}. \end{array} \right.$$

cannot learn the induction heads task for any A, B, C, Δ . Though they have not proved this.

- Their formulation was inspired from hypernetworks, gating, data-dependent transforms research BUT is not an GLU activation!

Mamba: Linear Time Sequence Modeling with Selective State Spaces

- Selection:

Table 11: (**Induction heads.**) Models are trained on sequence length $2^8 = 256$, and tested on various sequence lengths of $2^6 = 64$ up to $2^{20} = 1048576$. ✓ denotes perfect generalization accuracy, while ✗ denotes out of memory.

MODEL	PARAMS	TEST ACCURACY (%) AT SEQUENCE LENGTH														
		2^6	2^7	2^8	2^9	2^{10}	2^{11}	2^{12}	2^{13}	2^{14}	2^{15}	2^{16}	2^{17}	2^{18}	2^{19}	2^{20}
MHA-Abs	137K	✓	99.6	100.0	58.6	26.6	18.8	9.8	10.9	7.8	✗	✗	✗	✗	✗	✗
MHA-RoPE	137K	✓	✓	100.0	83.6	31.3	18.4	8.6	9.0	5.5	✗	✗	✗	✗	✗	✗
MHA-xPos	137K	✓	✓	100.0	99.6	67.6	25.4	7.0	9.0	7.8	✗	✗	✗	✗	✗	✗
H3	153K	✓	✓	100.0	80.9	39.5	23.8	14.8	8.2	5.9	6.6	8.2	4.7	8.2	6.3	7.4
Hyena	69M*	97.7	✓	100.0	✓	44.1	12.5	6.6	5.1	7.0	5.9	6.6	6.6	5.9	6.3	9.8
Mamba	74K	✓	✓	100.0	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓

* Most of the parameters are in learnable positional encodings.

Mamba: Linear Time Sequence Modeling with Selective State Spaces

Algorithm 1 SSM (S4)

Input: $x : (B, L, D)$
Output: $y : (B, L, D)$

- 1: $A : (D, N) \leftarrow \text{Parameter}$
 ▷ Represents structured $N \times N$ matrix
- 2: $B : (D, N) \leftarrow \text{Parameter}$
- 3: $C : (D, N) \leftarrow \text{Parameter}$
- 4: $\Delta : (D) \leftarrow \tau_\Delta(\text{Parameter})$
- 5: $\bar{A}, \bar{B} : (D, N) \leftarrow \text{discretize}(\Delta, A, B)$
- 6: $y \leftarrow \text{SSM}(\bar{A}, \bar{B}, C)(x)$
 ▷ Time-invariant: recurrence or convolution
- 7: **return** y

Algorithm 2 SSM + Selection (S6)

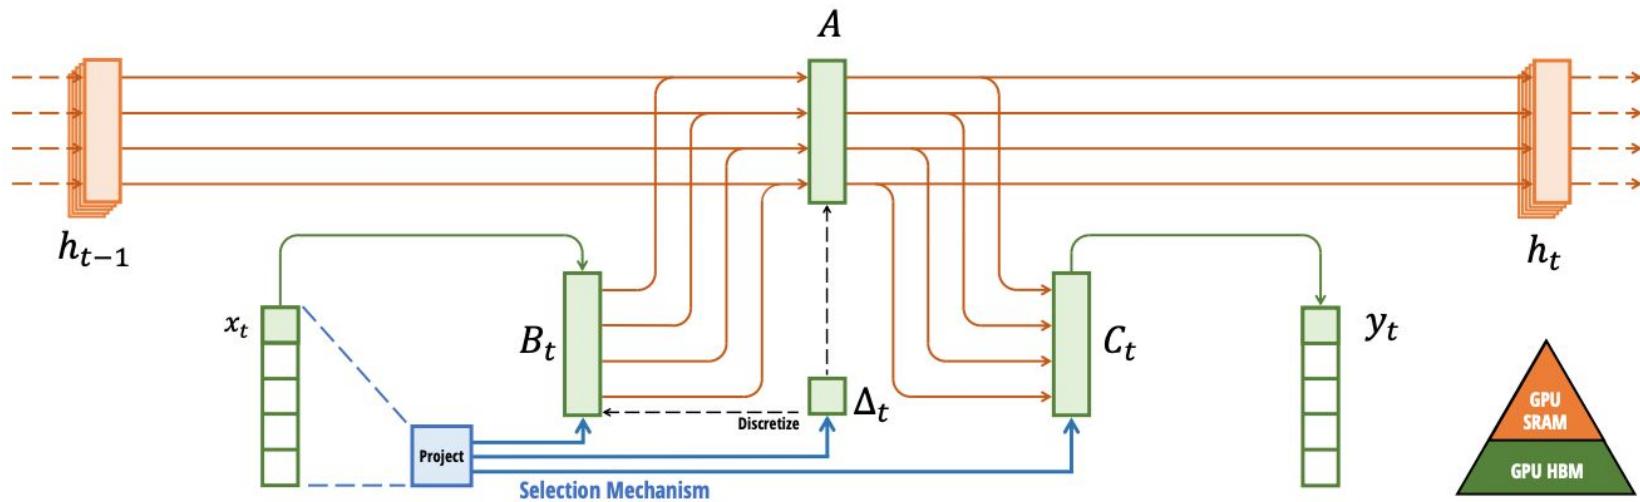
Input: $x : (B, L, D)$
Output: $y : (B, L, D)$

- 1: $A : (D, N) \leftarrow \text{Parameter}$
 ▷ Represents structured $N \times N$ matrix
- 2: $B : (B, L, N) \leftarrow s_B(x)$
- 3: $C : (B, L, N) \leftarrow s_C(x)$
- 4: $\Delta : (B, L, D) \leftarrow \tau_\Delta(\text{Parameter} + s_\Delta(x))$
- 5: $\bar{A}, \bar{B} : (B, L, D, N) \leftarrow \text{discretize}(\Delta, A, B)$
- 6: $y \leftarrow \text{SSM}(\bar{A}, \bar{B}, C)(x)$
 ▷ Time-varying: recurrence (*scan*) only
- 7: **return** y

Mamba: Linear Time Sequence Modeling with Selective State Spaces [the somewhat complicated part]

- The algorithm is **theoretically faster than S4**, for small state dimensions, but has a major problem.
 - Why faster than S4? Convolution is $O(B*L*D*\log(L))$ ($L*\log(L)$ because FFT), and Mamba conv is $O(B*L*D*N)$.
 - However, a naïve implementation would materialize the hidden state of dimensions $B*L*D*N$ in GPU HBM (High Bandwidth Memory).
 - Basically, CUDA kernels from Nvidia are not sufficient, so they write their own CUDA kernels to have **kernel fusion**.
 - This seems like fancy terms but are really easy concepts in reality. (easy to think of, not easy to implement!)

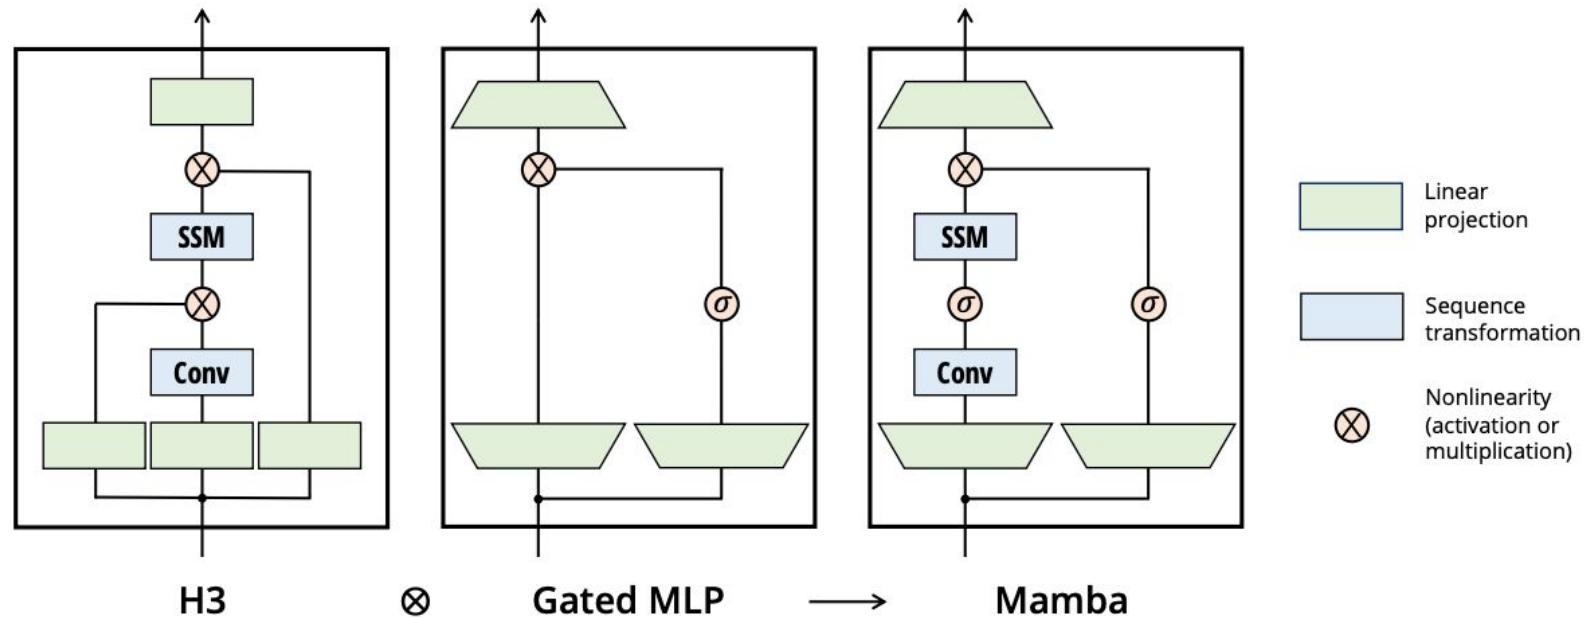
Mamba: Linear Time Sequence Modeling with Selective State Spaces



Mamba: Linear Time Sequence Modeling with Selective State Spaces

- The paper has 3 major contributions:
 - The LTI-drop (**selection** mechanism), ✓
 - Hardware-aware algorithm, ✓
 - Scaling ?
- Finally, this is the first paper on **this** branch of SSMs that scales their model to a few billion parameters to test on Language Modelling, and other real world tasks.
- Interestingly, they got rejected from ICLR'24 because at the time of submission they did not include LRA tasks.

Mamba: Linear Time Sequence Modeling with Selective State Spaces



Mamba: Linear Time Sequence Modeling with Selective State Spaces

- An important theorem.

Theorem 1. When $N = 1$, $\mathbf{A} = -1$, $\mathbf{B} = 1$, $s_\Delta = \text{Linear}(x)$, and $\tau_\Delta = \text{softplus}$, then the selective SSM recurrence (Algorithm 2) takes the form

$$\begin{aligned} g_t &= \sigma(\text{Linear}(x_t)) \\ h_t &= (1 - g_t)h_{t-1} + g_tx_t. \end{aligned} \tag{5}$$

Mamba: Linear Time Sequence Modeling with Selective State Spaces

- Some more discussion on the arch.
 - Variable Spacing / Filtering Context,
 - Boundary Resetting,
 - Interpretation of A, B, C and Δ .
- Major takeaway, Δ dictates a lot!

$\Delta \rightarrow \infty$ Resets **context**, massive focus on input.

$\Delta \rightarrow 0$ Ignores **input**, massive focus on context.

Mamba: Linear Time Sequence Modeling with Selective State Spaces

😢 Results (synthetic),

MODEL	ARCH.	LAYER	Acc.
S4	No gate	S4	18.3
-	No gate	S6	97.0
H3	H3	S4	57.0
Hyena	H3	Hyena	30.1
-	H3	S6	99.7
-	Mamba	S4	56.4
-	Mamba	Hyena	28.4
Mamba	Mamba	S6	99.8

Table 1: (**Selective Copying**)
Accuracy for combinations of architectures
and inner sequence layers.

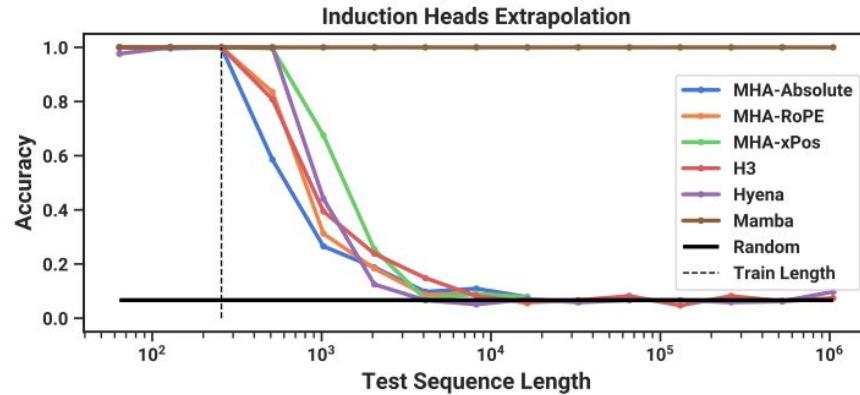
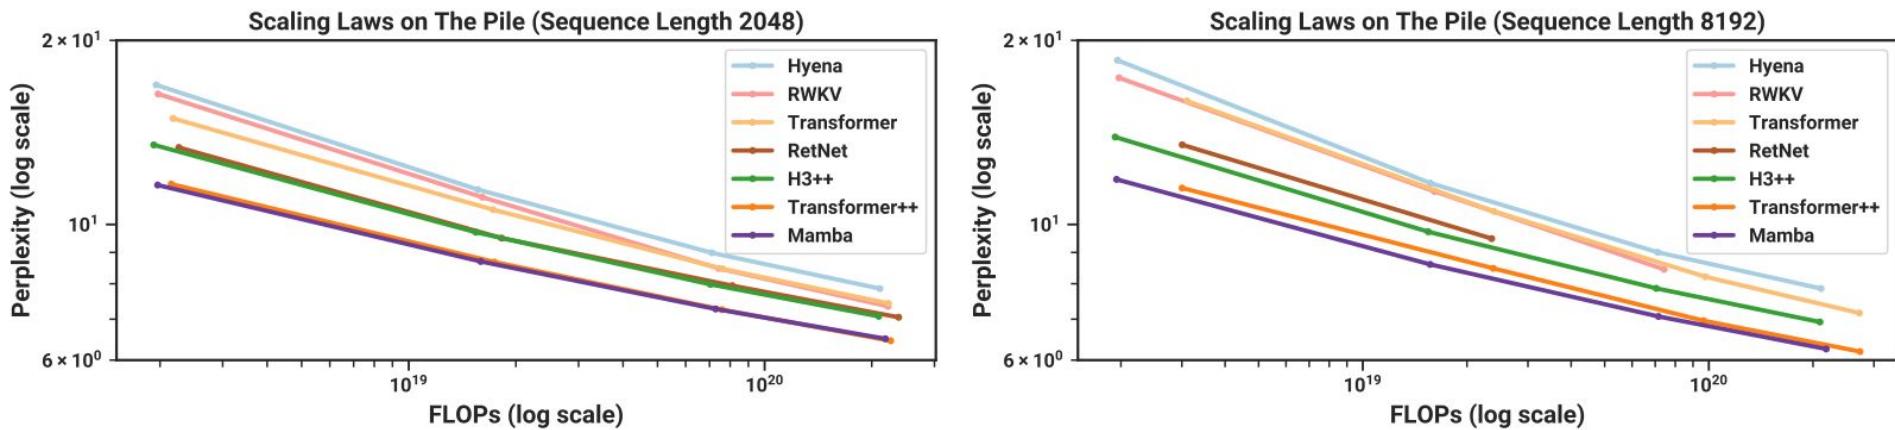


Table 2: (**Induction Heads**.) Models are trained on sequence length $2^8 = 256$, and tested on increasing sequence lengths of $2^6 = 64$ up to $2^{20} = 1048576$. Full numbers in Table 11.

Mamba: Linear Time Sequence Modeling with Selective State Spaces

😢 Results (scaling laws),



Mamba: Linear Time Sequence Modeling with Selective State Spaces

 Results
(language
modelling),

MODEL	TOKEN.	PILE PPL ↓	LAMBADA PPL ↓	LAMBADA ACC ↑	HELLASWAG ACC ↑	PIQA ACC ↑	Arc-E ACC ↑	Arc-C ACC ↑	WINOGRANDE ACC ↑	AVERAGE ACC ↑
Hybrid H3-130M	GPT2	—	89.48	25.77	31.7	64.2	44.4	24.2	50.6	40.1
Pythia-160M	NeoX	29.64	38.10	33.0	30.2	61.4	43.2	24.1	51.9	40.6
Mamba-130M	NeoX	10.56	16.07	44.3	35.3	64.5	48.0	24.3	51.9	44.7
Hybrid H3-360M	GPT2	—	12.58	48.0	41.5	68.1	51.4	24.7	54.1	48.0
Pythia-410M	NeoX	9.95	10.84	51.4	40.6	66.9	52.1	24.6	53.8	48.2
Mamba-370M	NeoX	8.28	8.14	55.6	46.5	69.5	55.1	28.0	55.3	50.0
Pythia-1B	NeoX	7.82	7.92	56.1	47.2	70.7	57.0	27.1	53.5	51.9
Mamba-790M	NeoX	7.33	6.02	62.7	55.1	72.1	61.2	29.5	56.1	57.1
GPT-Neo 1.3B	GPT2	—	7.50	57.2	48.9	71.1	56.2	25.9	54.9	52.4
Hybrid H3-1.3B	GPT2	—	11.25	49.6	52.6	71.3	59.2	28.1	56.9	53.0
OPT-1.3B	OPT	—	6.64	58.0	53.7	72.4	56.7	29.6	59.5	55.0
Pythia-1.4B	NeoX	7.51	6.08	61.7	52.1	71.0	60.5	28.5	57.2	55.2
RWKV-1.5B	NeoX	7.70	7.04	56.4	52.5	72.4	60.5	29.4	54.6	54.3
Mamba-1.4B	NeoX	6.80	5.04	64.9	59.1	74.2	65.5	32.8	61.5	59.7
GPT-Neo 2.7B	GPT2	—	5.63	62.2	55.8	72.1	61.1	30.2	57.6	56.5
Hybrid H3-2.7B	GPT2	—	7.92	55.7	59.7	73.3	65.6	32.3	61.4	58.0
OPT-2.7B	OPT	—	5.12	63.6	60.6	74.8	60.8	31.3	61.0	58.7
Pythia-2.8B	NeoX	6.73	5.04	64.7	59.3	74.0	64.1	32.9	59.7	59.1
RWKV-3B	NeoX	7.00	5.24	63.9	59.6	73.7	67.8	33.1	59.6	59.6
Mamba-2.8B	NeoX	6.22	4.23	69.2	66.1	75.2	69.7	36.3	63.5	63.3
GPT-J-6B	GPT2	—	4.10	68.3	66.3	75.4	67.0	36.6	64.1	63.0
OPT-6.7B	OPT	—	4.25	67.7	67.2	76.3	65.6	34.9	65.5	62.9
Pythia-6.9B	NeoX	6.51	4.45	67.1	64.0	75.2	67.3	35.5	61.3	61.7
RWKV-7.4B	NeoX	6.31	4.38	67.2	65.5	76.1	67.8	37.5	61.0	62.5

Thank you! Questions?